

Aggregate Programming

Foundations and Tools

Mirko Viroli

ALMA MATER STUDIORUM—Università di Bologna, Italy
`mirko.viroli@unibo.it`

Seminar at Università degli Studi di Urbino "Carlo Bo"
Corso di Laurea in Informatica Applicata
Urbino, 17/5/2016



The IoT is becoming a crowded and complex place..

Future and emerging Internet-of-things are witnessing..

- increasing availability of wearable / mobile / embedded / flying devices
 - increasing availability of heterogeneous wireless connectivity
 - increasing availability of computational resources (device/edge/cloud)
 - increasing production/analysis of data, everywhere, anytime
- ⇒ business / security / privacy concerns will probably be drivers, too



The IoT is becoming a crowded and complex place..

A plethora of programming models for “mobile/IoT applications”

- client side
 - ▶ single-device program: objects + functions + concurrency..
..threads/actors/futures/tasks/activities
 - ▶ device-centric interactions/protocols: using local APIs for
MoM/SOA/ad-hoc-communications
- server side
 - ▶ same interactions/protocols: MoM/SOA/ad-hoc-communications
 - ▶ storage by DB: OO, relational, NoSQL
 - ▶ coordination (orchestration, mediation, rules enactment)
 - ▶ situation recognition (online/offline, mining, business intelligence, stream processing)
- scalability in the server calls for cloudification
 - ▶ not really orthogonal to the whole programming model
 - ▶ it often dramatically affects system design



Implications

Where programming effort ends up?

- programs of clients and servers highly depend on
 - ▶ the chosen platform / API / communication technology
 - ▶ the number, type, and dislocation of involved devices
- ⇒ IoT systems tend to be very rigid, hard and costly to debug/maintain
- ⇒ design and deployments hardly tolerate changes

The technological result

- systems can't scale with complexity of behaviour
- very few of the opportunities of large-scale IoT are currently taken
 - ▶ virtually any computational mechanism (sensing, actuation, processing, storage)..
 - ▶ ..could involve spontaneous, adaptive cooperation of large sets of devices!
- how many large-scale deployments of adaptive IoT systems around?
- where are the Collective Adaptive Systems?

Implications

Where programming effort ends up?

- programs of clients and servers highly depend on
 - ▶ the chosen platform / API / communication technology
 - ▶ the number, type, and displacement of involved devices
- ⇒ IoT systems tend to be very rigid, hard and costly to debug/maintain
- ⇒ design and deployments hardly tolerate changes

The technological result

- systems can't scale with complexity of behaviour
- very few of the opportunities of large-scale IoT are currently taken
 - ▶ virtually any computational mechanism (sensing, actuation, processing, storage)..
 - ▶ ..could involve spontaneous, adaptive cooperation of large sets of devices!
- how many large-scale deployments of adaptive IoT systems around?
- where are the Collective Adaptive Systems?

What to do? A programming model perspective..

What do we lack in large-scale IoT systems?

- the plain old platform-independent programming abstraction
 - ⇒ fully grounding system design like objects did well.. in the past
 - ▶ delegating to the underlying platform virtually *all* deployment issues
 - ▶ automagically addressing non-functional issues (resilience, self-*)

The challenge

Just directly consider the worst scenario possible..

- zillion devices unpredictably moving in the environment
 - heterogeneous displacement, pervasive sensing/actuation
 - abstracting away from the possible multi-layered “server system”
 - ▶ whether with have fog++/cloud++ in background
- ⇒ but be ready to exploit the opportunities it creates!



What to do? A programming model perspective..

What do we lack in large-scale IoT systems?

- the plain old platform-independent programming abstraction
 - ⇒ fully grounding system design like objects did well.. in the past
 - ▶ delegating to the underlying platform virtually *all* deployment issues
 - ▶ automagically addressing non-functional issues (resilience, self-*)

The challenge

Just directly consider the worst scenario possible..

- zillion devices unpredictably moving in the environment
 - heterogeneous displacement, pervasive sensing/actuation
 - abstracting away from the possible multi-layered “server system”
 - ▶ whether with have fog++/cloud++ in background
- ⇒ but be ready to exploit the opportunities it creates!

What to do? A programming model perspective..

What do we lack in large-scale IoT systems?

- the plain old platform-independent programming abstraction
 - ⇒ fully grounding system design like objects did well.. in the past
 - ▶ delegating to the underlying platform virtually *all* deployment issues
 - ▶ automagically addressing non-functional issues (resilience, self-*)

The challenge

Just directly consider the worst scenario possible..

- zillion devices unpredictably moving in the environment
- heterogeneous displacement, pervasive sensing/actuation
- abstracting away from the possible multi-layered “server system”
 - ▶ whether with have fog++/cloud++ in background
 - ⇒ but be ready to exploit the opportunities it creates!

Let's try to program *that* “computational system”!

Abstract of the talk

Systems of interest: collective adaptive situated systems CASS

- (possibly very large scale) collective adaptive systems
- deployed in physical space (situated), i.e., IoT-oriented
- complex (open, dynamic, in need of much self-*)

Aggregate Computing

- The “good” computing/programming model for CASS
- It gives nice abstractions, promoting solid engineering principles
- Simple idea, few constructs, rather tractable, somehow *different*

This talk

1. Motivation and idea of aggregate computing
2. Some semi-technicalities and overview of results
3. State of toolchain

Abstract of the talk

Systems of interest: collective adaptive situated systems CASS

- (possibly very large scale) collective adaptive systems
- deployed in physical space (situated), i.e., IoT-oriented
- complex (open, dynamic, in need of much self-*)

Aggregate Computing

- The “good” computing/programming model for CASS
- It gives nice abstractions, promoting solid engineering principles
- Simple idea, few constructs, rather tractable, somehow *different*

This talk

1. Motivation and idea of aggregate computing
2. Some semi-technicalities and overview of results
3. State of toolchain

Abstract of the talk

Systems of interest: collective adaptive situated systems CASS

- (possibly very large scale) collective adaptive systems
- deployed in physical space (situated), i.e., IoT-oriented
- complex (open, dynamic, in need of much self-*)

Aggregate Computing

- The “good” computing/programming model for CASS
- It gives nice abstractions, promoting solid engineering principles
- Simple idea, few constructs, rather tractable, somehow *different*

This talk

1. Motivation and idea of aggregate computing
2. Some semi-technicalities and overview of results
3. State of toolchain

1 Aggregate Computing

2 Field Calculus

3 Tools



1 Aggregate Computing

2 Field Calculus

3 Tools



Let's follow standard “software engineering” process

Requirements and Analysis

- The customer does not mention “servers” or “connectivity”
- Different services to be implemented
- All in need of robustness at different levels
- Several common problems, to be “factored out”

Architectural design

- Depict strategies and abstractions in a platform-independent way
- Using concepts very near to the problem domain
- Identify common patterns

Detailed design and other stages

- Choose technologies, write APIs and component interfaces
- Implementation, Testing, Deployment

Let's follow standard “software engineering” process

Requirements and Analysis

- The customer does not mention “servers” or “connectivity”
- Different services to be implemented
- All in need of robustness at different levels
- Several common problems, to be “factored out”

Architectural design

- Depict strategies and abstractions in a platform-independent way
- Using concepts very near to the problem domain
- Identify common patterns

Detailed design and other stages

- Choose technologies, write APIs and component interfaces
- Implementation, Testing, Deployment

Let's follow standard “software engineering” process

Requirements and Analysis

- The customer does not mention “servers” or “connectivity”
- Different services to be implemented
- All in need of robustness at different levels
- Several common problems, to be “factored out”

Architectural design

- Depict strategies and abstractions in a platform-independent way
- Using concepts very near to the problem domain
- Identify common patterns

Detailed design and other stages

- Choose technologies, write APIs and component interfaces
- Implementation, Testing, Deployment

Broad research challenges

Computational/programming model for these services

- Programming as: “describing the problem, not hacking the solution!”
- Hiding complexity and resiliency “under-the-hood”
- How computation carries on is hidden as well, and intrinsically self-*

Grounding an effective tool-chain

- languages, compilers, simulators, scalable execution platforms

Supporting solid engineering principles

- checking/enacting functional/non-functional correctness
- supporting reuse of patterns, substitutability, compositionality

Chasing the true issue

- we should **fully** escape the single “device” abstraction

Broad research challenges

Computational/programming model for these services

- Programming as: “describing the problem, not hacking the solution!”
- Hiding complexity and resiliency “under-the-hood”
- How computation carries on is hidden as well, and intrinsically self-*

Grounding an effective tool-chain

- languages, compilers, simulators, scalable execution platforms

Supporting solid engineering principles

- checking/enacting functional/non-functional correctness
- supporting reuse of patterns, substitutability, compositionality

Chasing the true issue

- we should **fully** escape the single “device” abstraction

Broad research challenges

Computational/programming model for these services

- Programming as: “describing the problem, not hacking the solution!”
- Hiding complexity and resiliency “under-the-hood”
- How computation carries on is hidden as well, and intrinsically self-*

Grounding an effective tool-chain

- languages, compilers, simulators, scalable execution platforms

Supporting solid engineering principles

- checking/enacting functional/non-functional correctness
- supporting reuse of patterns, substitutability, compositionality

Chasing the true issue

- we should **fully** escape the single “device” abstraction

Approaches to “group interaction in space”

Survey of past approaches [Beal et.al., 2013]

- *Device abstractions* – make interaction implicit
NetLogo, Hood, TOTA, Gro, MPI, and the SAPERE approach
- *Pattern languages* – supporting composability of spatial behaviour
Growing Point, Origami Shape, various selforg pattern langs
- *Information movement* – gathering in space, moving elsewhere
TinyDB and Regiment
- *Foundation* – giving linguistic means for group interactions in space
 3π , Shape Calculus, bi-graphs, KLAIM, $\sigma\tau$ -linda, SCEL
- *Spatial computing* – program space-time behaviour of systems
Proto, MGS

Our approach

- Combining the above efforts of “macro” programming
- Taking some of those ideas to the extreme consequences

Approaches to “group interaction in space”

Survey of past approaches [Beal et.al., 2013]

- *Device abstractions* – make interaction implicit
NetLogo, Hood, TOTA, Gro, MPI, and the SAPERE approach
- *Pattern languages* – supporting composability of spatial behaviour
Growing Point, Origami Shape, various selforg pattern langs
- *Information movement* – gathering in space, moving elsewhere
TinyDB and Regiment
- *Foundation* – giving linguistic means for group interactions in space
 3π , Shape Calculus, bi-graphs, KLAIM, $\sigma\tau$ -linda, SCEL
- *Spatial computing* – program space-time behaviour of systems
Proto, MGS


Our approach

- Combining the above efforts of “macro” programming
- Taking some of those ideas to the extreme consequences

Manifesto of aggregate computing

Motto: program the aggregate, not individual devices!

1. The reference computing machine
⇒ an aggregate of devices as single “body”, fading to the actual *space*
2. The reference elaboration process
⇒ atomic manipulation of a collective data structure (a *field*)
3. The actual networked computation
⇒ a proximity-based self-org system hidden “under-the-hood”



1 Aggregate Computing

2 Field Calculus


3 Tools




Computational Fields [Mamei et.al., 2009, Beal et.al., 2013]

Traditionally a map: $Space \mapsto Values$


- possibly: evolving over time, dynamically injected, stabilising
- smoothly adapting to very heterogeneous domains
- more easily “understood” on continuous and flat spatial domains
- ranging to: booleans, reals, vectors, functions



boolean channel in 2D



numeric partition in 2D



real-valued gradient in 3D




(Computational) Fields revisited

A field as a space-time structure: $\phi : D \mapsto V$

- *event* E : a triple $\langle \delta, t, p \rangle$ – device δ , “firing” at time t in position p
- *events domain* D : a coherent set of events (devices cannot move too fast)
- *field values* V : any data value

Domain




(Computational) Fields revisited


A field as a space-time structure: $\phi : D \mapsto V$

- *event* E : a triple $\langle \delta, t, p \rangle$ – device δ , “firing” at time t in position p
- *events domain* D : a coherent set of events (devices cannot move too fast)
- *field values* V : any data value

Domain



Field




(Computational) Fields revisited


A field as a space-time structure: $\phi : D \mapsto V$

- *event* E : a triple $\langle \delta, t, p \rangle$ – device δ , “firing” at time t in position p
- *events domain* D : a coherent set of events (devices cannot move too fast)
- *field values* V : any data value

Domain



Field




will later show only snapshots of fields in 2D space..



The “channel” example: computing a redundant route

How would you program it?




how could a program be platform-independent,
unaware of global map, resilient to changes, faults...



The “channel” example: computing a redundant route

How would you program it?




how could a program be platform-independent,
unaware of global map, resilient to changes, faults,...



Aggregate programming as a functional approach

Functionally composing fields

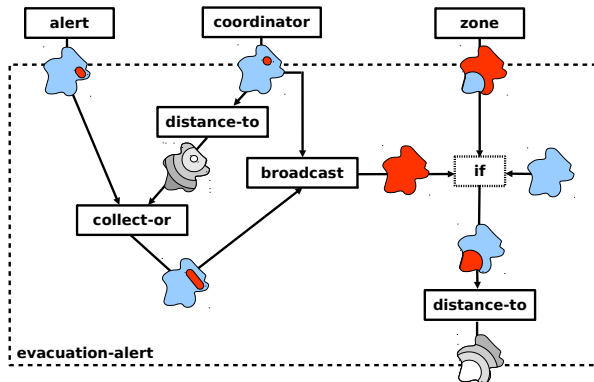
- Inputs: sensor fields, Output: actuator field
 - Computation is a pure function over fields (time embeds state!)
- ⇒ for this to be practical/expressive we need a good programming language



Crowd evacuation as a field computation

Computing by purely functional composition of space-time fields

- Inputs: sensor fields, Output: actuator field
 - Computation is a pure function over fields (time embeds state!)
- ⇒ for this to be practical/expressive we need a good programming language



Field calculus [Damiani & Viroli & Beal & Pianini, FORTE2015]

Key idea

- a sort of λ -calculus with “everything is a field” philosophy!

Syntax (slightly refactored, semi-formal version of FORTE's)

$e ::= x \mid v \mid e(e_1, \dots, e_n) \mid \text{rep}(e_0)\{e\} \mid \text{nbr}\{e\}$	(expr)
$v ::= < \text{standard-values} > \mid \lambda$	(value)
$\lambda ::= f \mid o \mid (\bar{x}) \Rightarrow e$	(functional value)
$F ::= \text{def } f(\bar{x}) \{e\}$	(function definition)

Few explanations





- v includes numbers, booleans, strings, ..
..tuples/vectors/maps/any-ADT (of expressions)
- f is a user-defined function
- o is a built-in functional operator (mostly pure math or a sensor)

Intuition of global-level semantics

The four main constructs at work

⇒ values, application, evolution, and interaction – in aggregate guise

• $e ::= \dots \mid v \mid e(e_1, \dots, e_n) \mid \text{rep}(e_0)\{e\} \mid \text{nbr}\{e\}$



A mini-tutorial: functions, repetitions, neighbouring

```
1: 1
2: 2 + 3
3: pair(10,20)
4: random()
5: sense(1)
6: sense(1) ? 10 : 20
7: mid()
8: minHood(nbrRange)
```

```
1: rep(0){ (x) => x + 1 }
2: rep(random()){ (x) => x }
3: rep(0){ (x) => x + rep(random()){ (y) => y } }
```

```
1: maxHood( nbr{ sense(1) } )
2: sumHood( nbr{ 1 } )
```

```
1: rep(0){ (x) => max( sense(1), maxHood( nbr{ x } ) ) }
2: rep(Infinity) { (d) => sense(1) ? 0 : minHood( nbr{d} + 1 ) }
3: rep(Infinity) { (d) => sense(1) ? 0 : minHood( nbr{d} + nbrRange ) }
```



A mini-tutorial: functions, repetitions, neighbouring

```
1: 1                ;; values become constant fields
2: 2 + 3            ;; math is done infix
3: pair(10,20)      ;; fst, snd to extract
4: random()         ;; note iterative execution..
5: sense(1)         ;; a boolean sensor
6: sense(1) ? 10 : 20 ;; muxing
7: mid()            ;; unique identifiers
8: minHood(nbrRange) ;; distance of closest neighbour
```

```
1: rep(0){ (x) => x + 1 } ;; counting the number of rounds
2: rep(random()){ (x) => x } ;; stable random
3: rep(0){ (x) => x + rep(random()){ (y) => y } } ;; counting at different velocities
```

```
1: maxHood( nbr{ sense(1) } ) ;; maximum value of sensor in neighbours
2: sumHood( nbr{ 1 } ) ;; number of neighbours
```

```
1: rep(0){ (x) => max( sense(1), maxHood( nbr{ x } ) ) } ;; gossiping max of sense(1)
2: rep(Infinity) { (d) => sense(1) ? 0 : minHood( nbr{d} + 1 ) } ;; hop-count
3: rep(Infinity) { (d) => sense(1) ? 0 : minHood( nbr{d} + nbrRange ) } ;; gr
```



Intuition of global-level semantics

Value v

- A field constant in space and time, mapping any event to v

Function application $e(e_1, \dots, e_n)$

- e evaluates to a field of functions, assume it ranges to $\lambda_1, \dots, \lambda_n$
- this naturally induces a partition of the domain D_1, \dots, D_n
- now, join the fields: $\forall i, \lambda_i(e_1, \dots, e_n)$ restricted in D_i


Repetition $\text{rep}(e_0)\{e_\lambda\}$

- the value of e_0 where the restricted domain “begins”
- elsewhere, unary function e_λ is applied to previous value at each device

Neighbouring field construction $\text{nbr}\{e\}$

- at each event gathers most recent value of e in neighbours (in restriction)
- ..what is neighbour is orthogonal (i.e., physical proximity)

The restriction trick: branching behaviour



if as a space-time branching construct

if(e-bool){e-then}**else**{e-else}

≈

(e-bool ? ()=>{e-then} : ()=>{e-else})()


More advanced patterns

- spread code, in different versions in different regions
- have different regions/device run different programs

Aggregate programming as a functional approach

Functionally composing fields

- ...so, is field calculus language practical/expressive?



The channel pattern

```
def gradient(source){ ;; reifying minimum distance from source
  rep(Infinity) { ;; distance is infinity initially
    (distance) => source ? 0 : minHood( nbr{distance} + nbrRange )
  } }

def distance(source, dest) { ;; propagates minimum distance between source and dest
  snd( ;; returning the second component of the pair
    rep(pair(Infinity, Infinity)) { ;; computing a field of pairs (distance,value)
      (distanceValue) => source ? pair(0, gradient(dest)) :
        minHood( ;; propagating as a gradient, using for first component of the pair
          pair(fst(nbr{distanceValue}) + nbrRange, snd(nbr{distanceValue})))
    } ) }

def dilate(region, width) { ;; a field of booleans
  gradient(region) < width
}

;; Here the "aggregate" nature of our approach gets revealed
def channel(source, dest, width) {
  dilate( gradient(source) + gradient(dest) <= distance(source,dest), width )
}
```



Builtin functions exploited


- **?:** — Java-like (though, call-by-value) ternary operator
- **nbrRange** — maps each device to a neighbour field of estimated distances
- **minHood** — in each device, collapse a neighbour field into its minimum value
- **sumHood** — in each device, collapse a neighbour field into sum of values
- ***, -, *, /, >, ...** — usual math, applied also pointwise to fields
- **pair, fst, snd** — construction/selection for pairs



Crowd evacuation as a field computation

Computing by purely functional composition of space-time fields

- Inputs: sensor fields, Output: actuator field
 - Computation is a pure function over fields (time embeds state!)
- ⇒ for this to be practical/expressive we need a good programming language



Evacuation example

```
def distance-to(source){ ;; reifying minimum distance from source
  rep(Infinity) { ;; distance is infinity initially
    (distance) => source ? 0 : minHood( nbr{distance} + nbrRange )
  } }

def broadcast(source, v) { ;; propagates minimum distance between source and dest
  snd( ;; returning the second component of the pair
    rep(pair(Infinity, v)) { ;; computing a field of pairs (distance,value)
      (distanceValue) => source ? pair(0, distance-to(v)) :
        minHood( ;; propagating as a gradient, using for first component of the pair
          pair(fst(nbr{distanceValue}) + nbrRange, snd(nbr{distanceValue})))
    } ) }

def collect-or (potential, value){;; Collects 'value' by descending 'potential', by 'or'
  rep(value){
    (v) => anyHood( nbr{find-parent potential()} = uid ? nbr{v} : false )
    or value
  } }

def evacuation-alert (zone, coordinator, alert){
  distance-to(
    if(zone){false} else {
      broadcast(coordinator,collect-or(distance-to(coordinator),alert))
    }
  )
}
```



On expressiveness of the field calculus

Practically, we can express:

- complex spreading / aggregation / decay functions
- spatial leader election, partitioning, consensus
- distributed spatio-temporal sensing and situation recognition
- dynamic deployment/spreading of code (via lambda)
- implicit/explicit device selection of what code execute
- “collective teams” forming based on the selected code



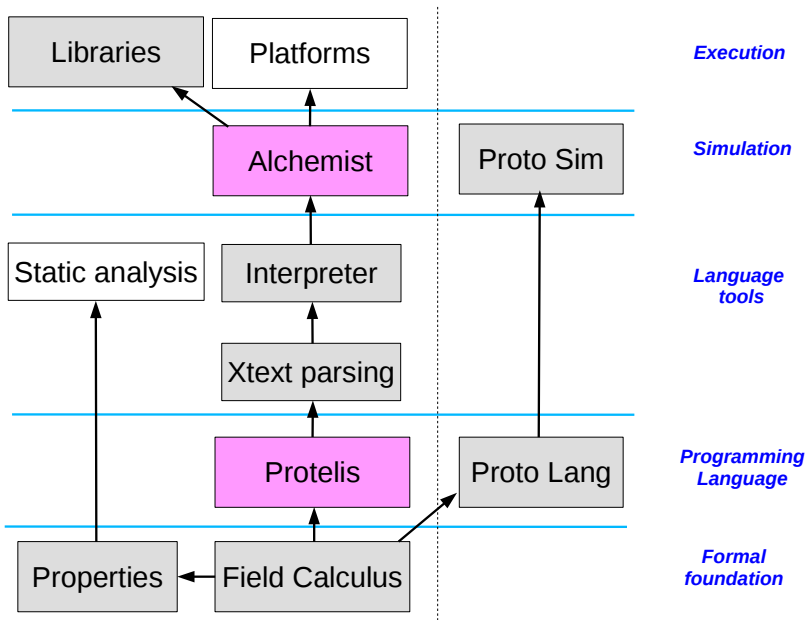
1 Aggregate Computing

2 Field Calculus

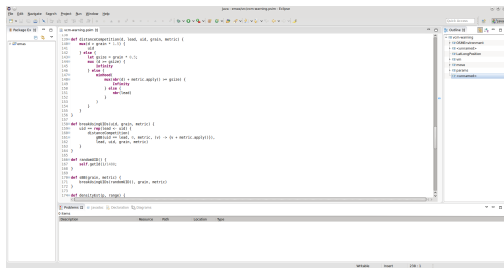
3 Tools



Current tool-chain for aggregate computing



Protelis + Alchemist [SAC-2015]



Protelis language: <http://protelis.org/>

- Field calculus in disguised and full-blown version
- Java-like syntax and Java API integration

Alchemist simulator: <http://alchemist.apice.unibo.it/>

- A general-purpose simulator with pluggable specification language
- XText/Eclipse integration
- Support from working with Maps, Traces, Paths, Movement models

Conclusions

Aggregate Computing

- a new paradigm for developing large-scale situated systems
- a bunch of results and tools emerged, many to come
- we're always eager to find new collaborations!

Acknowledgments

- Jacob Beal (BBN, USA)
- Ferruccio Damiani (UNITO)
- Danilo Pianini (UNIBO)



References I

- [IEEE Computer 48(9), 2015] Beal, J., Pianini, D. and Viroli, M. (2015).
Aggregate programming for the Internet of Things.
IEEE Computer, 48(9) 2015.
- [Beal et.al., 2013] Beal, J., Dulman, S., Usbeck, K., Viroli, M., and Correll, N. (2013).
Organizing the aggregate: Languages for spatial computing.
In Mernik, M., editor, *Formal and Practical Aspects of Domain-Specific Languages: Recent Developments*, chapter 16, pages 436–501. IGI Global.
A longer version available at: <http://arxiv.org/abs/1202.5509>.
- [SCW-2014] Beal, J., Viroli, M., and Damiani, F. (2014).
Towards a unified model of spatial computing.
In *7th Spatial Computing Workshop (SCW 2014)*, AAMAS 2014, Paris, France.
- [Fruin, 1971] Fruin, J. (1971).
Pedestrian Planning and Design.
Metropolitan Association of Urban Designers and Environmental Planners.
- [Mamei et.al., 2009] Mamei, M. and Zambonelli, F. (2009).
Programming pervasive and mobile computing applications: The tota approach.
ACM Transactions on Software Engineering and Methodologies, 18(4).
- [SAC-2015] Pianini, D., Beal, J., and Viroli, M. (2015).
Practical aggregate programming with PROTELIS.
In *ACM Symposium on Applied Computing (SAC 2015)*.
To appear.



References II

- [Damiani & Viroli & Beal & Pianini, FORTE2015] Damiani, F., Viroli, M., Pianini, D., and Beal, J. (2015).
Code mobility meets self-organisation: a higher-order calculus of computational fields.
In *Formal Techniques for Distributed Objects, Components, and Systems*, volume 9039 of *LNCS*, pages 113–128. Springer.
- [SASO-2015] Viroli, M., Beal, J., Damiani, F. and Pianini, D. (2015).
Efficient Engineering of Complex Self-Organising Systems by Self-Stabilising Fields
In *IEEE Conference on Self-Adaptive and Self-Organising Systems*.
- [Beal, SAC2009] Beal, j. (2009).
Flexible self-healing gradients
In *ACM Symposium on Applied Computing*, pp. 1197–1201.

